

Composing candidates buys more accuracy per second than selecting one

LibriSpeech test-other. Each curve walks the candidate pool from $K = 1$ to 32 (Whisfusion) or 64 (Drax); cost is the measured decode plus encode time per utterance on one T4.

